



MODULE 5

Module V (8 Hours)

The Memory System – basic concepts, semiconductor RAM memories - organization – static and dynamic RAM, Structure of larger memories, semiconductor ROM memories, Speed, Size and cost, Cache memory – mapping functions – replacement algorithms, **Virtual memory – paging and segmentation.**

- Hamacher, Vranesic & Zaky, “Computer Organization” (5th Ed), McGraw Hill.

VIRTUAL MEMORY

- In most modern computer systems, the **physical main memory is not as large as the address space of the processor.**
 - For example, a processor that issues 32-bit addresses has an addressable space of 4G bytes. The size of the main memory in a typical computer with a 32-bit processor may range from 1G to 4G bytes.
- If a program does not completely fit into the main memory, the parts of it **not currently being executed are stored on a secondary storage device**, typically a magnetic disk.
- As these parts are needed for execution, they must **first be brought into the main memory**, possibly replacing other parts that are already in the memory.
- These actions are performed **automatically** by the operating system, using a scheme known as **virtual memory**.

VIRTUAL MEMORY

- Application programmers **need not be aware** of the limitations imposed by the available main memory.
- Under a virtual memory system, programs and hence the processor, reference instructions and data in an address space that is **independent of the available physical main memory space**.
- The binary addresses that the processor issues for either instructions or data are called **virtual or logical addresses**.
- These addresses are **translated into physical addresses** by a combination of hardware and software components.
- If a virtual address refers to a part of the program or data space that is **currently in the physical memory**, then the contents of the appropriate location in the main memory are accessed immediately.
- Otherwise, the contents of the referenced address **must be brought into a suitable location in the memory** before they can be used.

VIRTUAL MEMORY

[Hamacher, Vranesic & Zaky, "Computer Organization" (5th Ed), McGraw Hill]

- Figure 15 shows a typical organization that implements **virtual memory**.
- A special hardware unit, called the **Memory Management Unit (MMU)**, keeps track of which parts of the virtual address space are in the physical memory.
- When the desired data or instructions are in the main memory, the **MMU translates the virtual address into the corresponding physical address**.
- Then, the requested memory access proceeds in the usual manner.

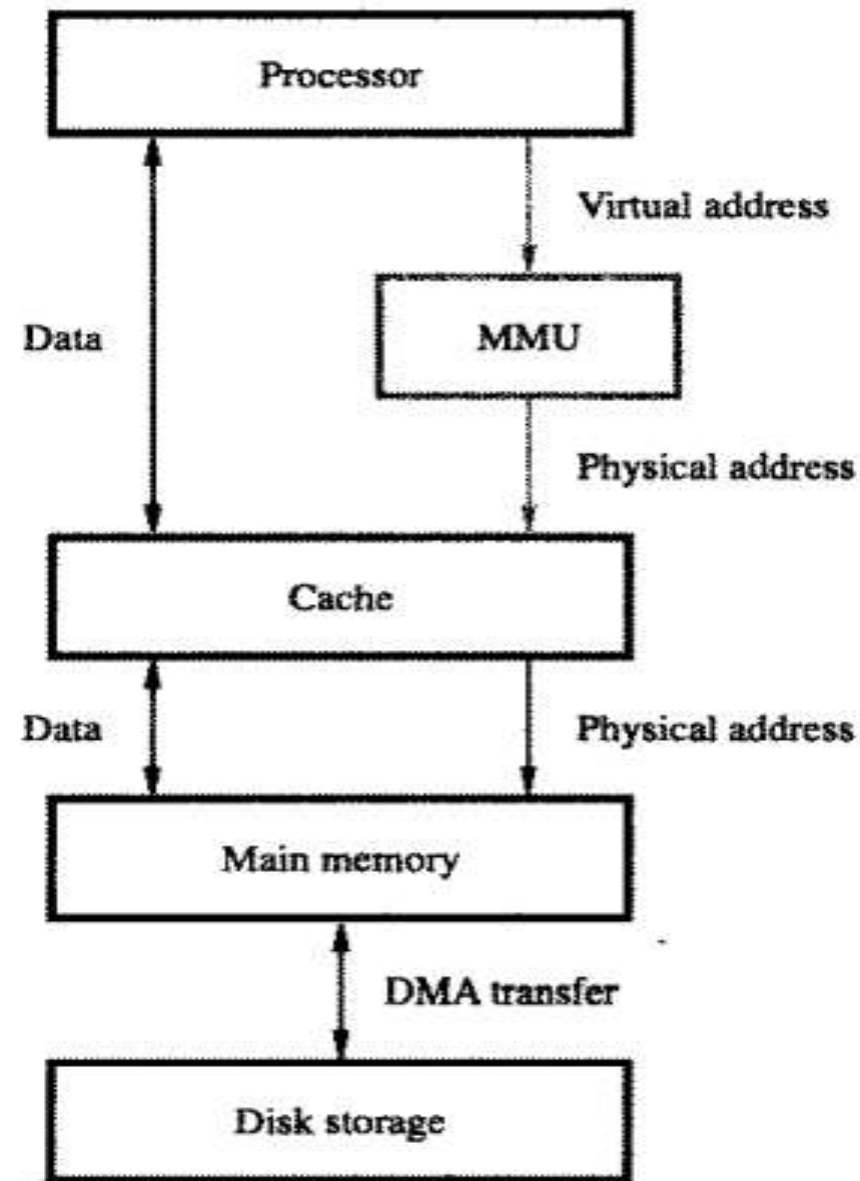


Figure 15 Virtual memory organization.

VIRTUAL MEMORY

[Hamacher, Vranesic & Zaky, "Computer Organization" (5th Ed), McGraw Hill]

- If the data are **not in the main memory**, the MMU causes the operating system to transfer the data from the disk to the memory.
- Such transfers are performed using the **DMA scheme**.




Figure 15 Virtual memory organization.

Address Translation

- A simple method for translating virtual addresses into physical addresses is to assume that all programs and data are composed of **fixed-length units** called **pages**, each of which consists of a **block of words** that occupy contiguous locations in the main memory.
- **Pages** commonly range from **2K to 16K bytes** in length.
- They constitute the **basic unit of information** that is transferred between the main memory and the disk whenever the MMU determines that a transfer is required.

Address Translation

- Pages should **not be too small**, because the access time of a magnetic disk is much longer (several milliseconds) than the access time of the main memory.
- The reason for this is that it takes a **considerable amount of time to locate the data on the disk**, but once located, the data can be transferred at a rate of several megabytes per second.
- On the other hand, if **pages are too large**, it is possible that a substantial portion of a page may not be used, yet this unnecessary data will occupy valuable space in the main memory.

Address Translation

- The **cache** bridges the **speed gap** between the processor and the main memory and is implemented in hardware.
- The **virtual-memory** mechanism bridges **the size and speed gaps** between the main memory and secondary storage and is usually implemented in part by software techniques.
- Conceptually, cache techniques and virtual-memory techniques are **very similar**. They differ mainly in the details of their implementation.

Address Translation

[Hamacher, Vranesic & Zaky, "Computer Organization" (5th Ed), McGraw Hill]

- A virtual-memory address-translation method based on the concept of **fixed-length pages** is shown schematically in Figure 16.
- Each **virtual address** generated by the processor, whether it is for an instruction fetch or an operand load/store operation, is interpreted as a **virtual page number** (high-order bits) followed by an **offset** (low-order bits) that specifies the location of a particular byte (or word) within a page.

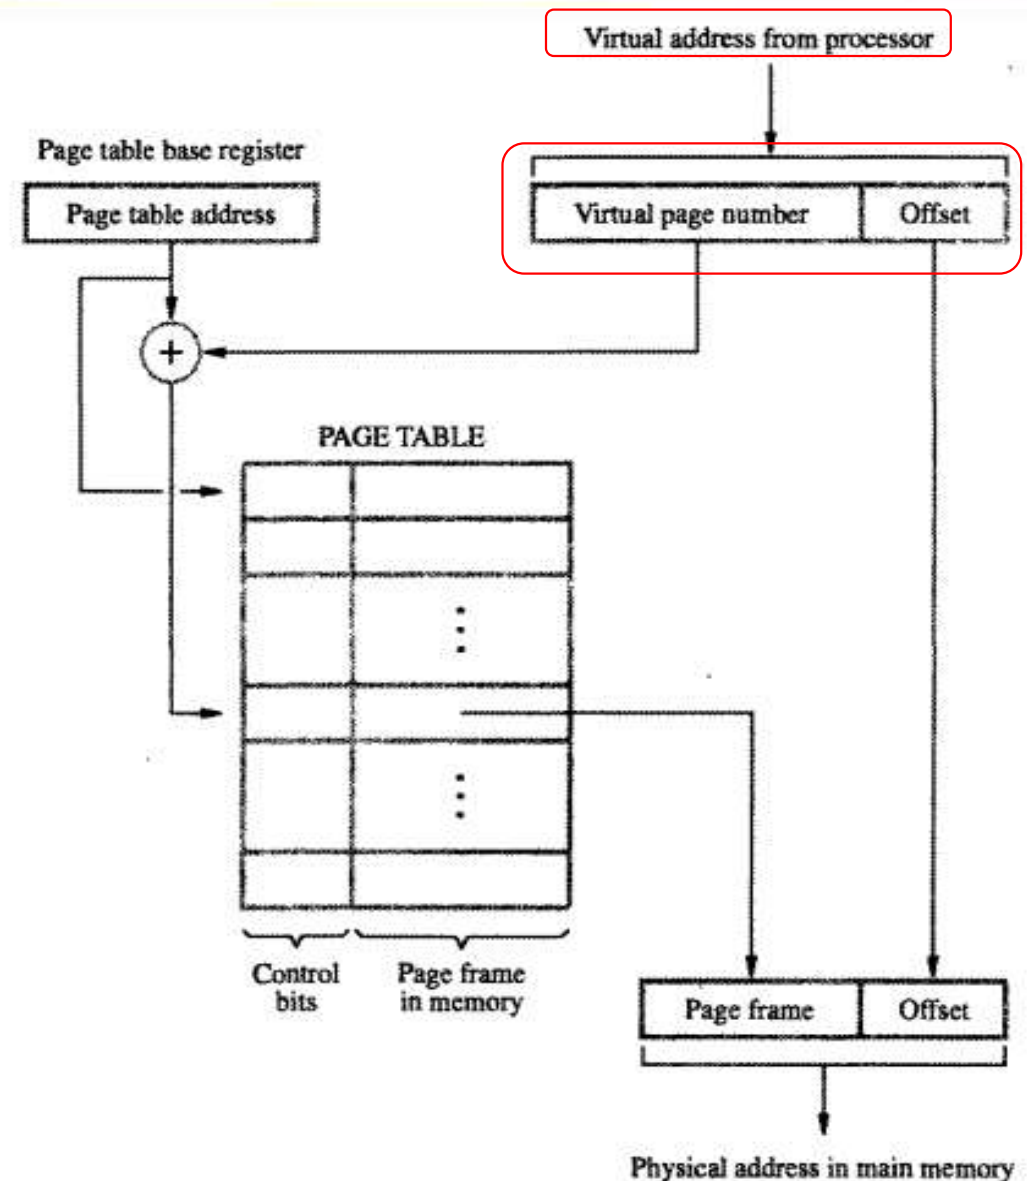


Figure 16 Virtual-memory address translation.

Address Translation

[Hamacher, Vranesic & Zaky, "Computer Organization" (5th Ed), McGraw Hill]

- Information about the **main memory location** of each page is kept in a **page table**.
- This information includes the **main memory address** where the page is stored and the current status of the page.
- An area in the main memory that can hold one page is called a **page frame**.

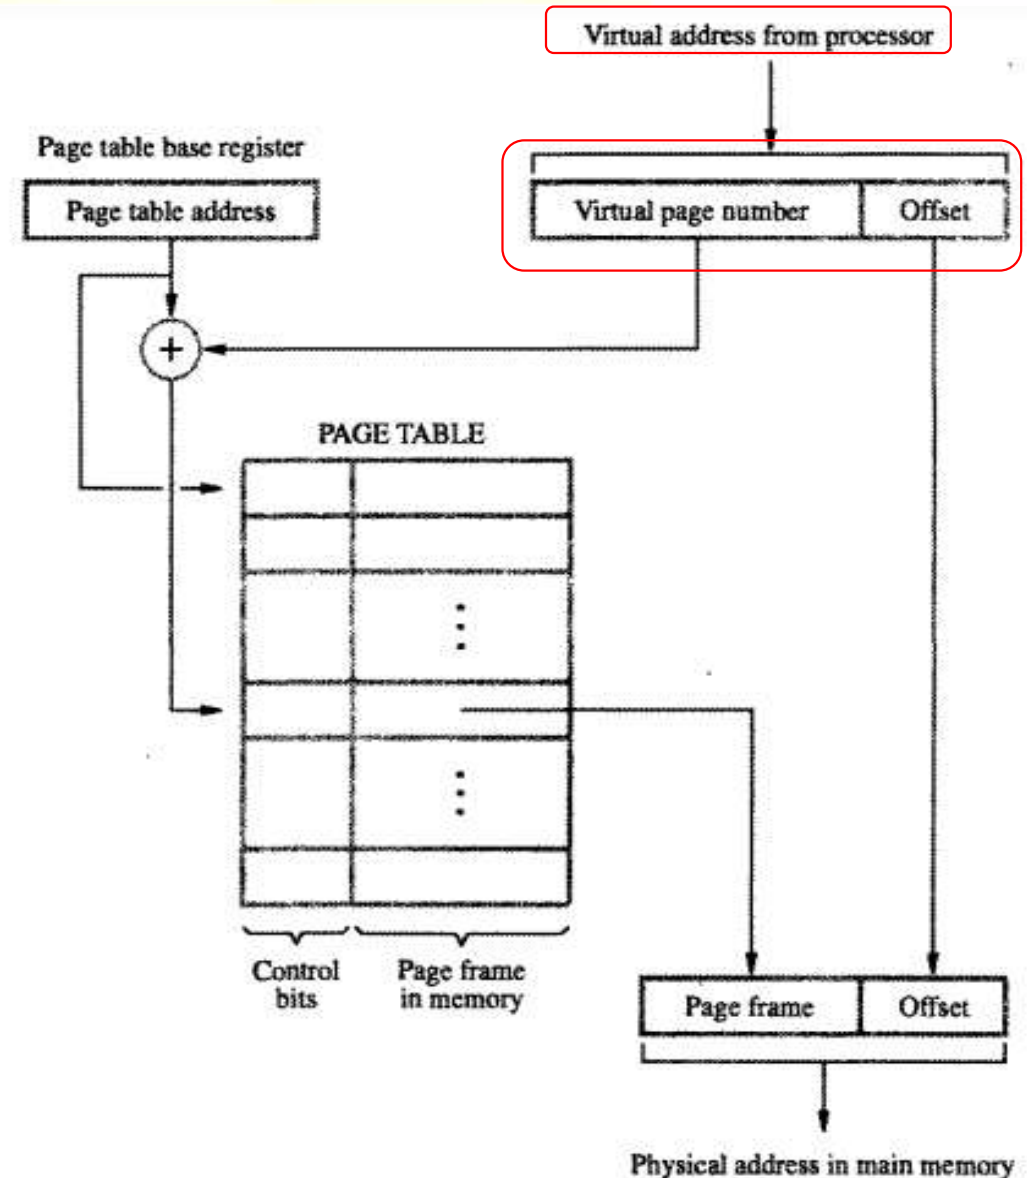


Figure 16 Virtual-memory address translation.

Address Translation

[Hamacher, Vranesic & Zaky, "Computer Organization" (5th Ed), McGraw Hill]

- The **starting address** of the page table is kept in a **page table base register**.
- By **adding** the virtual page number to the contents of this register, the **address of the corresponding entry** in the page table is obtained.
- The contents of this location give the **starting address of the page** if that page currently resides in the main memory.




Figure 16 Virtual-memory address translation.

Address Translation

[Hamacher, Vranesic & Zaky, "Computer Organization" (5th Ed), McGraw Hill]

- Each entry in the page table also includes **control bits** that describe the **status** of the page while it is in the main memory.
- One bit indicates the **validity of the page**, that is, whether the page is **actually loaded in the main memory**.
- It allows the operating system to **invalidate the page** without actually removing it.

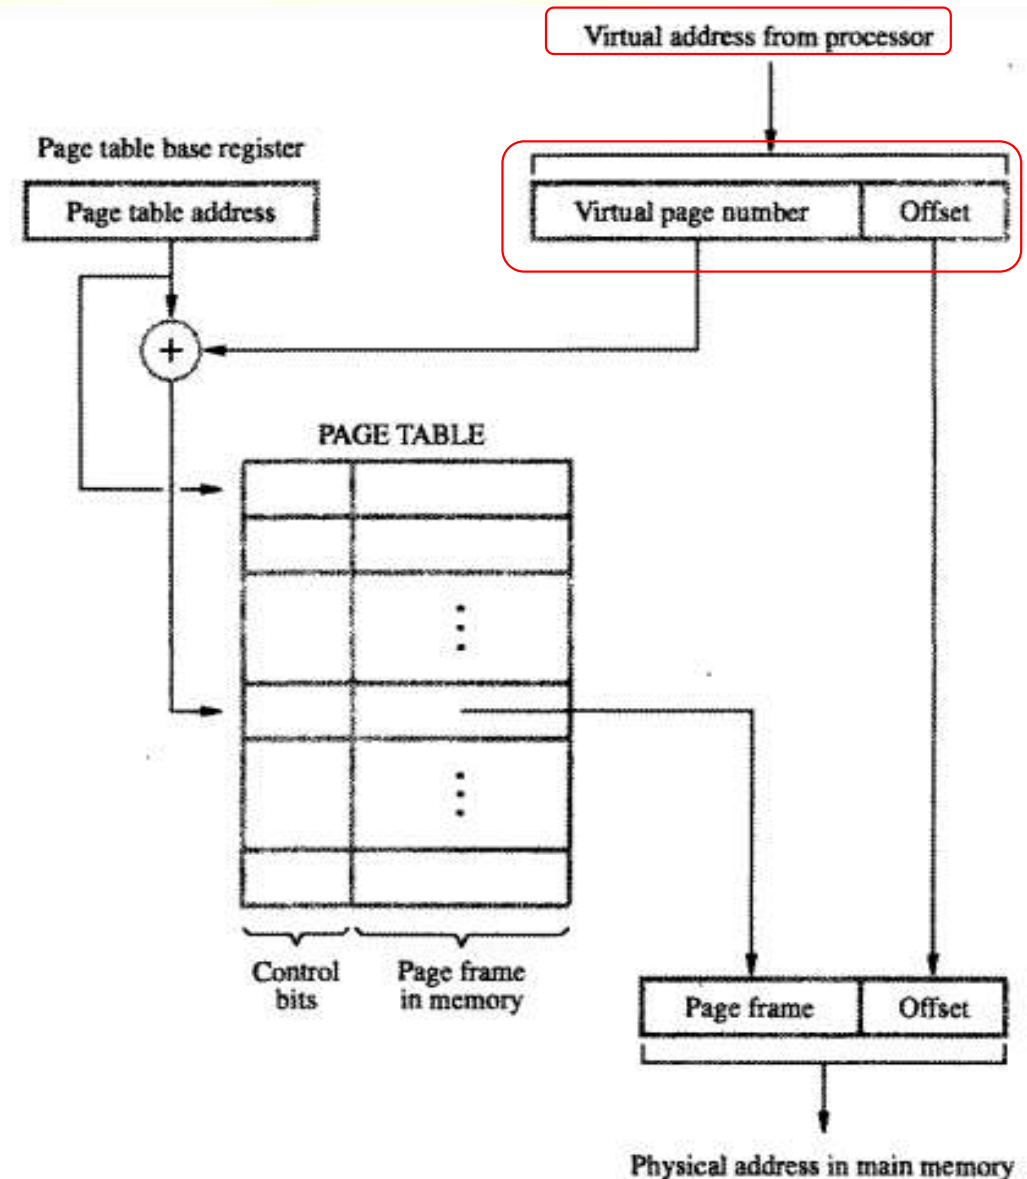


Figure 16 Virtual-memory address translation.

Address Translation

[Hamacher, Vranesic & Zaky, "Computer Organization" (5th Ed), McGraw Hill]

- Another **bit** indicates whether the page has been **modified** during its residency in the memory.
- As in cache memories, this information is needed to determine whether the page should be **written back to the disk** before it is removed from the main memory to make room for another page.
- Other **control bits** indicate various **restrictions** that may be imposed on accessing the page.
- For example, a program may be given **full read and write permission**, or it may be restricted to **read accesses only**.




Figure 16 Virtual-memory address translation.

Translation Lookaside Buffer

- The **page table information** is used by the **MMU** for every read and write access.
- Ideally, the page table should be situated within the **MMU**.
- Unfortunately, the **page table may be rather large** and since the MMU is normally implemented as part of the processor chip, it is impossible to include the complete table within the MMU.
- Instead, a copy of **only a small portion of the table is accommodated within the MMU**, and the **complete table is kept in the main memory**.

Translation Lookaside Buffer

- The portion maintained within the MMU consists of the entries corresponding to the **most recently accessed pages**.
- They are stored in a small table, usually called the **Translation Lookaside Buffer (TLB)**.
- The TLB functions as a **cache for the page table in the main memory**.
- Each entry in the TLB includes **a copy of the information in the corresponding entry in the page table**.
- In addition, it includes the **virtual address of the page**, which is needed to search the TLB for a particular page.

Translation Lookaside Buffer

- Figure 17 shows a possible organization of a TLB that uses the **associative mapping** technique.
- **Set-associative mapped TLBs** are also found in commercial products.
- The contents of the TLB be **coherent** with the contents of page tables in the memory.

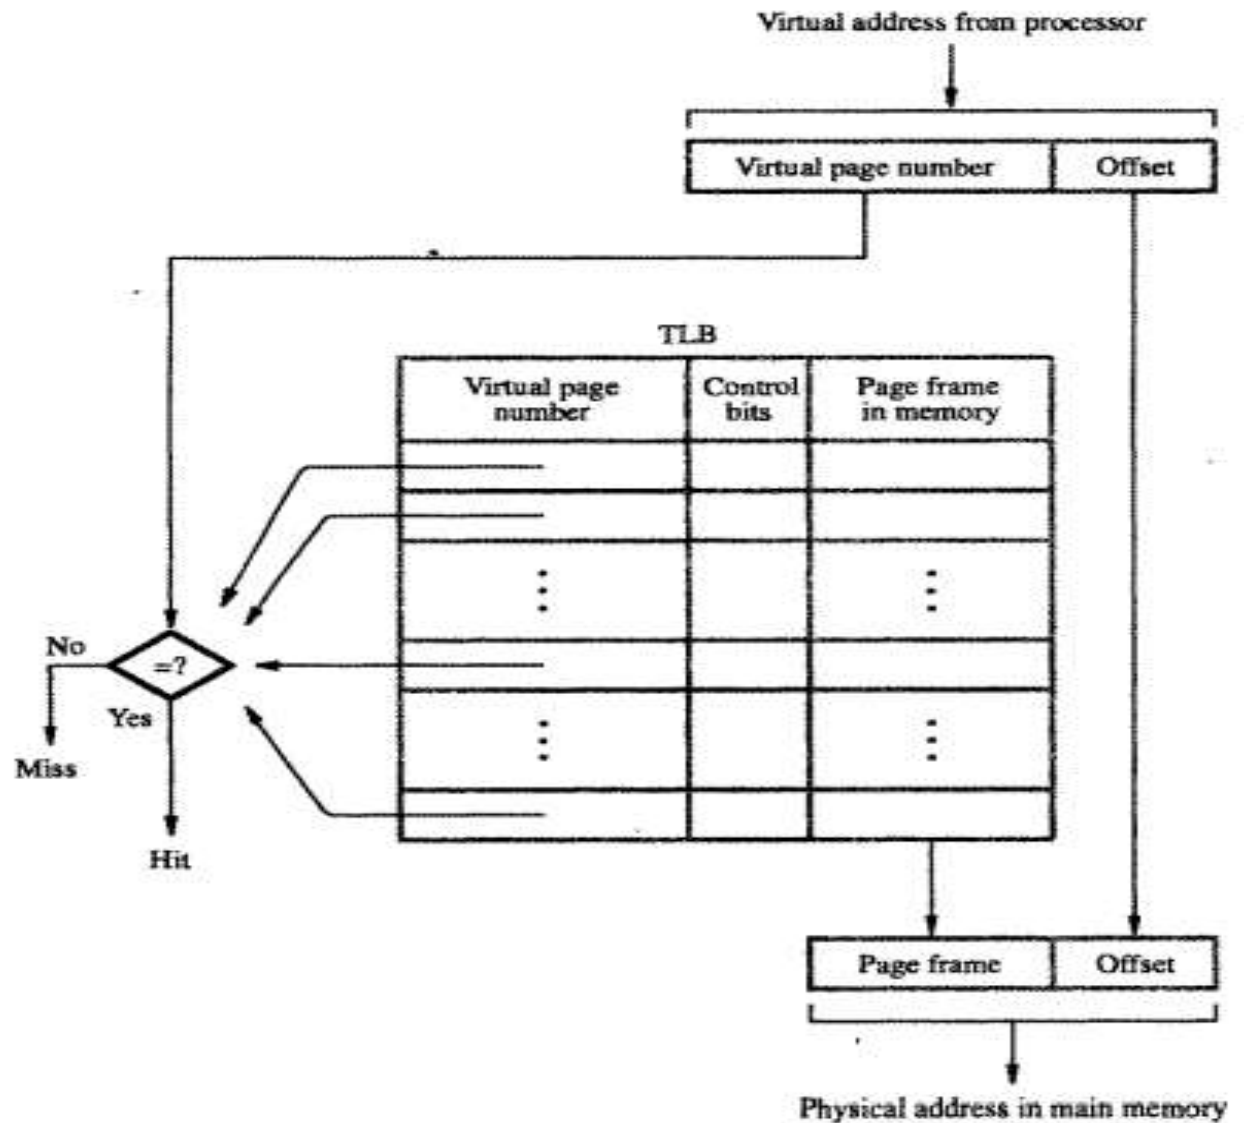


Figure 17 Use of an associative-mapped TLB.

Translation Lookaside Buffer

- It is essential to ensure that the contents of the TLB are always the **same as** the contents of page tables in the memory.
- When the operating system changes the contents of a page table, it must **simultaneously invalidate the corresponding entries in the TLB**.
- One of the **control bits** in the TLB is provided for this purpose.
- When an entry is **invalidated**, the TLB acquires the **new information from the page table** in the memory as part of the MMU's normal response to access misses.

Translation Lookaside Buffer

Address translation proceeds as follows.

- Given a virtual address, the MMU looks in the TLB for the referenced page.
- If the page table entry for this page is found in the TLB, the physical address is obtained immediately.
- If there is a miss in the TLB, then the required entry is obtained from the page table in the main memory and the TLB is updated.

Page Faults

- When a program generates an access request to a page that is not in the main memory, a **page fault** is said to have occurred.
- The **entire page** must be brought from the disk into the memory before access can proceed.
- When it detects a **page fault**, the MMU asks the operating system to intervene by **raising an exception (interrupt)**.
- Processing of the program that generated the page fault is **interrupted**, and control is transferred to the operating system.
- The operating system **copies the requested page** from the disk into the main memory.
- Since this process involves **a long delay**, the operating system may begin execution of another program whose pages are in the main memory.
- When **page transfer is completed**, the execution of the interrupted program is **resumed**.

Page Faults

- When the MMU raises an interrupt to indicate a page fault, the instruction that requested the memory access may have been **partially executed**.
- It is essential to ensure that the interrupted program **continues correctly** when it resumes execution.
- There are **two options**. Either the execution of the interrupted instruction continues from **the point of interruption**, or the instruction must be restarted.
- The **design of a particular processor** dictates which of these two options is used.

Page Faults

- If a new page is brought from the disk when the main memory is **full**, it must **replace** one of the resident pages.
- The problem of choosing which page to remove is just as **critical** here as it is in a cache and the observation that programs spend most of their time in a **few localized areas** also applies.
- Because main memories are considerably larger than cache memories, it should be **possible to keep relatively larger portions of a program in the main memory**.
- This **reduces the frequency** of transfers to and from the disk.

Page Faults

- Concepts similar to the **LRU replacement algorithm** can be applied to page replacement and the **control bits** in the page table entries can be used to **record usage history**.
- One simple scheme is based on a control bit that is **set to 1** whenever the corresponding page is referenced (accessed).
- The operating system **periodically clears this bit** in all page table entries, thus providing a simple way of determining which pages **have not been used recently**.

Page Faults

- A modified page has to be **written back to the disk** before it is removed from the main memory.
- It is important to note that the **write-through protocol**, which is useful in the framework of cache memories, is **not suitable for virtual memory**.
- The **access time of the disk is so long** that it does not make sense to access it frequently to write small amounts of data.

Page Faults

- Looking up entries in the TLB introduces some delay, slowing down the operation of the MMU.
- Here again we can take advantage of the property of locality of reference.
- It is likely that many successive TLB translations involve addresses on the same program page. This is particularly likely when fetching instructions.
- Thus, address translation time can be reduced by keeping the most recently used TLB entries in a few special registers that can be accessed quickly.