

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

DNS, or the Domain Name System, **translates human readable domain names** (for example, `www.amazon.com`) **to machine readable IP addresses** (for example, `192.0.2.44`).

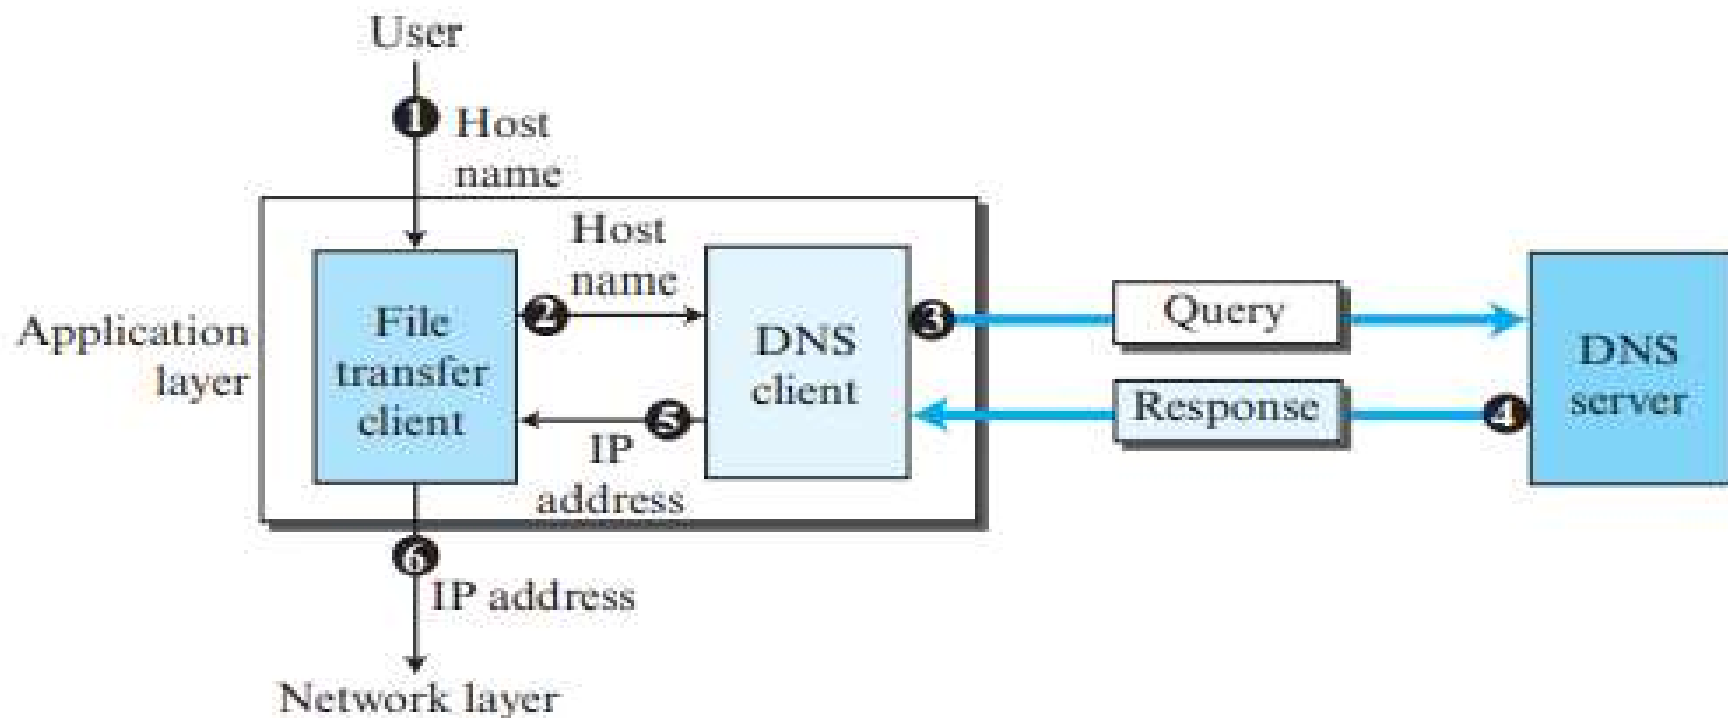
- Since the Internet is so huge today, a **central directory** system cannot hold all the mapping. In addition, if the central computer fails, the whole communication network will collapse.
- A better solution is to **distribute** the information among many computers in the world.
- In this method, the host that needs mapping can contact the closest computer holding the needed information. This method is used by the **Domain Name System (DNS)**.

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

- Figure 1.50 shows how TCP/IP uses a DNS client and a DNS server to map a name to an address.

Figure 1.50 Purpose of DNS



STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

The following **six steps** map the host name to an IP address:

1. The user passes the **host name** to the file transfer client.
2. The file transfer client passes the **host name** to the DNS client.
3. Each computer, after being booted, knows the address of one DNS server. The **DNS client sends a message to a DNS server with a query** that gives the file transfer server name using the known **IP address of the DNS server**.
4. The DNS server responds with the IP address of the **desired file transfer server**.
5. The DNS client **passes the IP address** to the file transfer server.
6. The file transfer client now uses the received **IP address** to access the file transfer server.

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Name Space

- A name space that **maps each address to a unique name** can be organized in two ways: flat or hierarchical.
- In a **flat name space**, a name is assigned to an address. A name in this space is a sequence of characters without structure.
- In a **hierarchical name space**, each name is made of several parts. The first part can define the **nature** of the organization, the second part can define the **name** of an organization, the third part can define **departments** in the organization, and so on.
- In this case, the authority to assign and control the name spaces can be **decentralized**.


STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Domain Name Space

- To have a **hierarchical name space**, a domain name space was designed.
- In this design the names are defined in an **inverted-tree structure** with the root at the top. The tree can have only 128 levels: level 0 (root) to level 127 (see Figure 1.51).

Figure 1.51 *Domain name space*



STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Label

- Each node in the tree has a label, which is a string with a maximum of 63 characters.
- The **root label** is a null string (empty string).
- DNS requires that children of a node (nodes that branch from the same node) have **different labels**, which guarantees the **uniqueness** of the domain names.

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

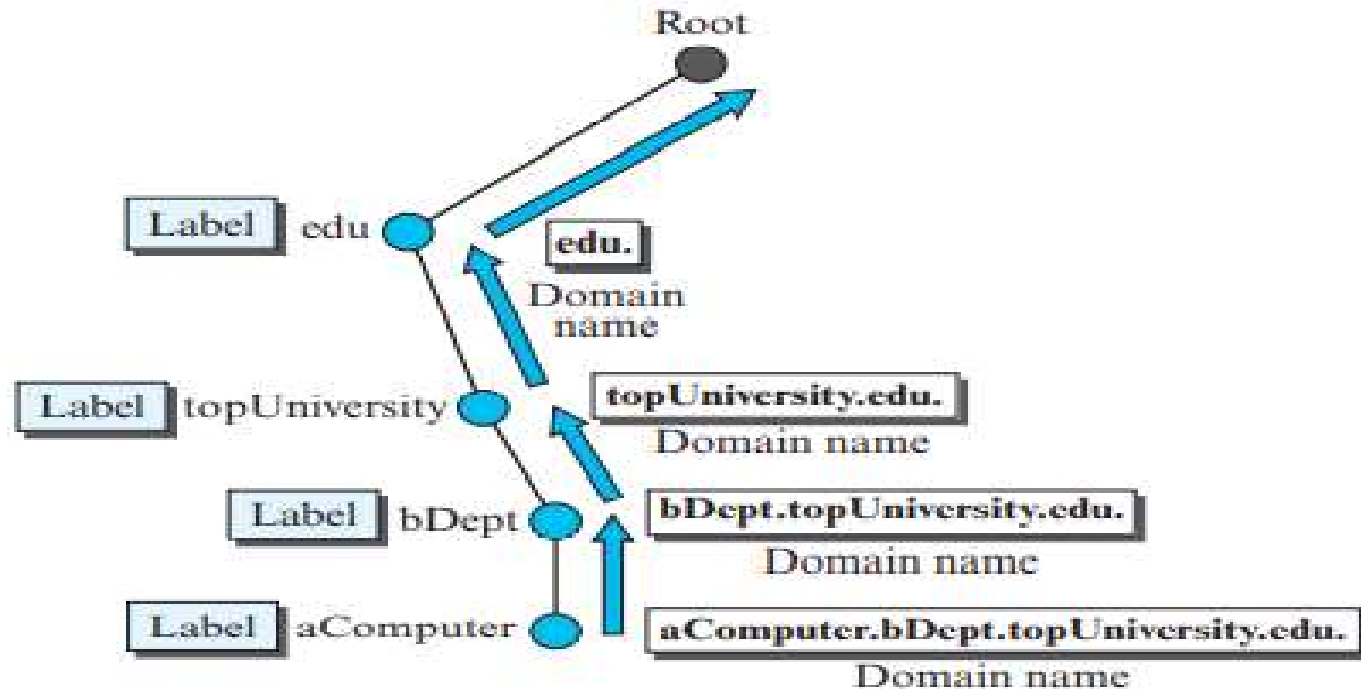
Domain Name

- Each node in the tree has a domain name.
- A full domain name is a **sequence of labels separated by dots (.)**.
- The domain names are always read **from the node up to the root**.
- The last label is the label of the **root** (null).
- This means that a full domain name always ends in a **null label**, which means the last character is a dot because the null string is nothing.
- Figure 1.52 shows some domain names.

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Figure 1.52 Domain names and labels



[Behrouz A Forouzan, Firouz Mosharraf, "Computer Networks: A top down Approach", McGraw Hill Education]

- If a label is **terminated by a null string**, it is called a **fully qualified domain name (FQDN)**.
- If a label is **not terminated by a null string**, it is called a **partially qualified domain name (PQDN)**.


STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Domain

- A domain is a **subtree** of the domain name space.
- The name of the domain is the name of the node at the top of the subtree.
- Note that a domain may itself be divided into domains.

Figure 1.53 Domains



STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Distribution of Name Space

- The information contained in the domain name space must be **stored**.
- It is very **inefficient** and also **not reliable** to have just one computer store such a huge amount of information.


Hierarchy of Name Servers

- The solution to these problems is to distribute the information among many computers called **DNS servers**.
- One way to do this is to divide the whole space into many **domains**.
- DNS allows domains to be divided further into smaller domains (subdomains).
- In other words, we have a **hierarchy of servers** in the same way that we have a hierarchy of names (see Figure 1.54).

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Figure 1.54 *Hierarchy of name servers*




STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Zone

- Since the complete **domain name** hierarchy cannot be stored on a single server, it is divided among many servers.
- What a server is responsible for or has authority over is called a **zone**.
- We can define a zone as a **contiguous part of the entire tree**.
- The server makes a database called a **zone file** and keeps all the information for every node under that domain.

Figure 1.55 Zone



STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Root Server

- A **root server** is a server whose zone consists of the **whole tree**.
- A root server usually does not store any information about domains but **delegates its authority** to other servers, keeping references to those servers.

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Primary and Secondary Servers

- DNS defines **two types** of servers: primary and secondary.
- A **primary server** is a server that stores a file about the zone for which it is **an authority**. It is **responsible for creating, maintaining, and updating the zone file**. It stores the zone file on a local disk.
- A **secondary server** is a server that transfers the complete information about a zone from another server (primary or secondary) and stores the file on its local disk.
- A **primary server** loads all information from the disk file; the **secondary server** loads all information from the primary server.

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

DNS in the Internet

- DNS is a protocol that can be used in different platforms.
- In the Internet, the domain name space (tree) was originally divided into three different sections: generic domains, country domains, and the inverse domain.

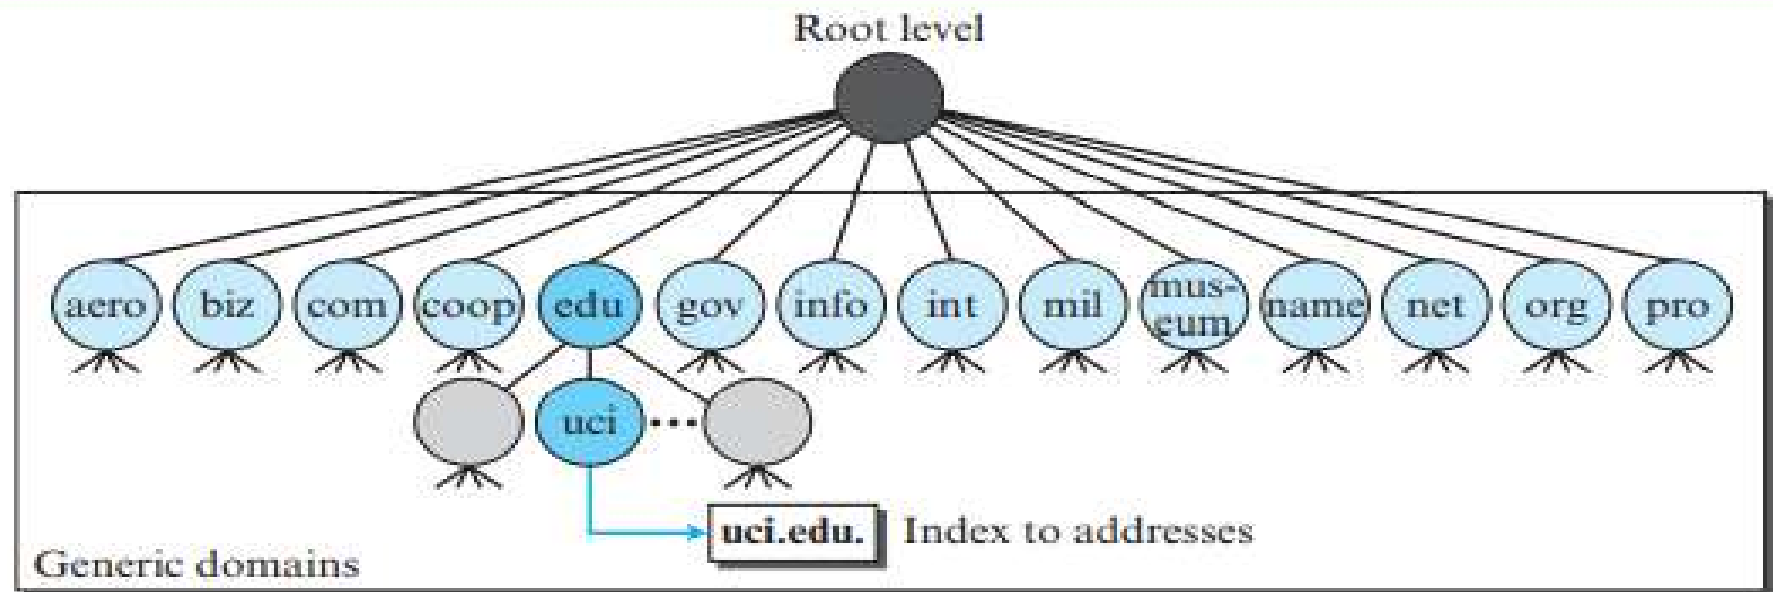
STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Generic Domains

- The **generic domains** define **registered hosts** according to their generic behavior.
- Each node in the tree defines a domain, which is an index to the domain name space database (see Figure 1.56).

Figure 1.56 Generic domains




STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Country Domains

- The country domains section uses **two-character country abbreviations** (e.g., us for United States, in for India).
- Second labels can be **organizational**, or they can be more specific, national designations.

Figure 1.57 Country domains



Irvine university in the state of California in the United States

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Resolution

- Mapping a name to an address is called **name-address resolution**.
- DNS is designed as a client-server application.
- A host that needs to **map an address to a name or a name to an address** calls a DNS client called a resolver.
 - Recursive Resolution
 - Iterative Resolution

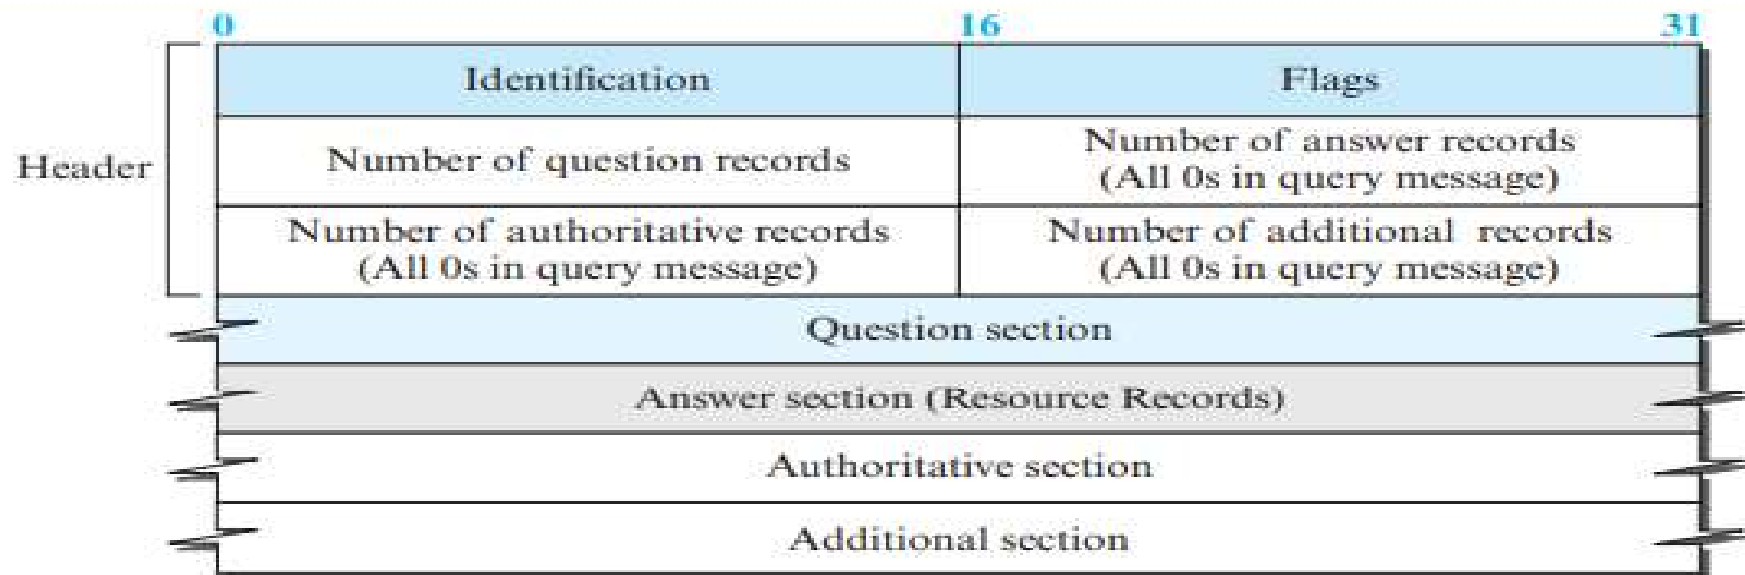
STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

DNS Messages

- To retrieve information about hosts, DNS uses two types of messages: **query** and **response**. Both types have the **same format** as shown in Figure 1.58.

Figure 1.58 *DNS message*



Note:

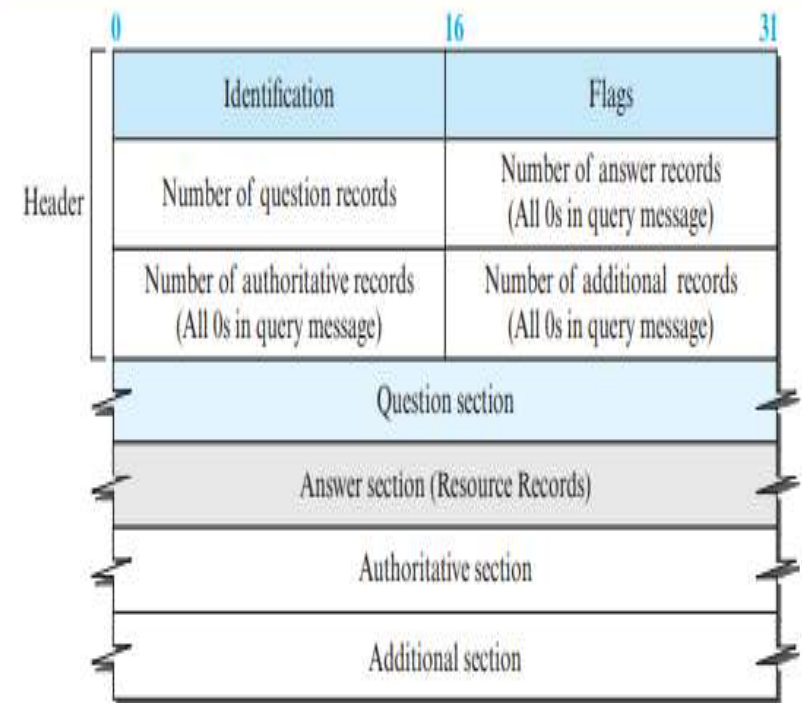
The query message contains only the question section.
The response message includes the question section,
the answer section, and possibly two other sections.

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

- The **identification field** is used by the client to match the response with the query.
- The **flag field** defines whether the message is a query or response. It also includes status of error.
- The next four fields in the header define the **number of each record type** in the message.
- The **question section**, which is included in the query and repeated in the response message, consists of one or more question records.

Figure 1.58 DNS message



Note:

The query message contains only the question section.
The response message includes the question section, the answer section, and possibly two other sections

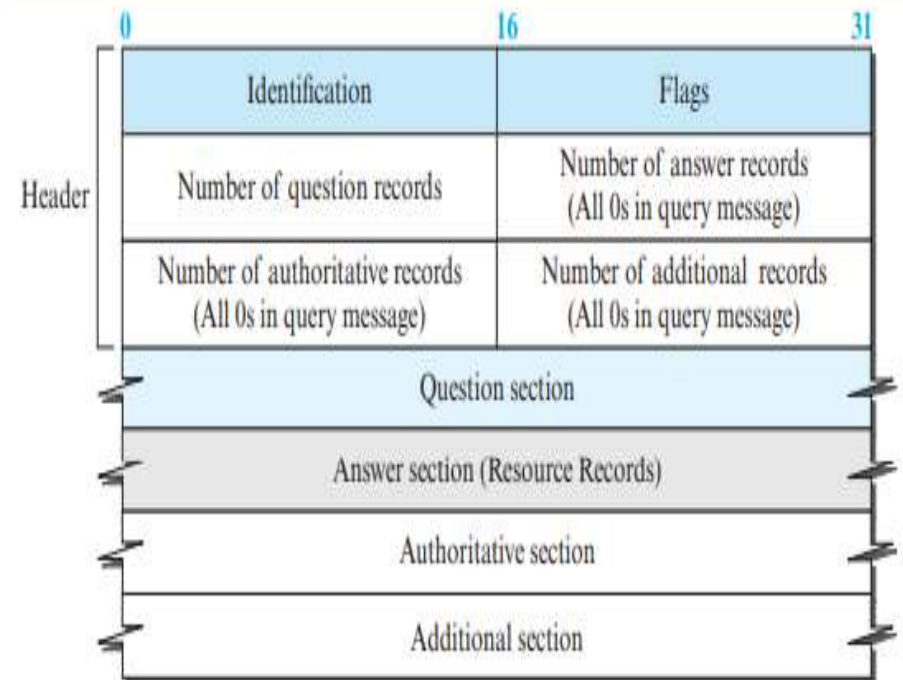
[Behrouz A Forouzan, Firouz Mosharraf, "Computer Networks: A top down Approach", McGraw Hill Education]

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

- The **answer section** consist of one or more resource records. It is present only in response messages.
- The **authoritative section** gives information (domain name) about one or more authoritative servers for the query.
- The **additional information** section provides additional information that may help the resolver.

Figure 1.58 DNS message



Note:

The query message contains only the question section.
The response message includes the question section, the answer section, and possibly two other sections.

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Encapsulation

- DNS can use either UDP or TCP.
- In both cases the well-known port used by the server is port 53.
- UDP is used when the size of the response message is less than 512 bytes because most UDP packages have a 512-byte packet size limit.
- If the size of the response message is more than 512 bytes, a TCP connection is used.

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

Registrars

How are new domains added to DNS?

- This is done through a **registrar**, a commercial entity accredited by ICANN(Internet Corporation for Assigned Names and Numbers).
- A registrar first verifies that the requested domain name is **unique** and then enters it into the DNS database. A fee is charged.

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

DDNS

- The **DNS master file must be updated** dynamically.
- The **Dynamic Domain Name System** (DDNS) was devised to respond to this need.
- In DDNS, when **a binding between a name and an address is determined**, the information is sent, usually by **DHCP** (Dynamic Host Configuration Protocol) to a primary DNS server.
- The primary server updates the zone. The secondary servers are notified either actively or passively.
- In **active notification**, the primary server sends a message to the secondary servers about the change in the zone, whereas in passive notification, the secondary servers periodically check for any changes.
- To provide security and prevent unauthorized changes in the DNS records, DDNS can use an **authentication mechanism**.

STANDARD CLIENT-SERVER APPLICATIONS

Domain Name System (DNS)

- To protect DNS, IETF has devised a technology named **DNS Security (DNSSEC)** that provides message origin authentication and message integrity using a security service called **digital signature**.

PEER-TO-PEER PARADIGM

P2P Networks


- We first need to divide the P2P networks into two categories: **centralized and decentralized**.

Centralized Networks

- In a centralized P2P network, **the directory system** – listing of the peers and what they offer – uses the client-server paradigm, but the **storing and downloading of the files are done using the peer-to-peer paradigm** (a hybrid P2P network).

PEER-TO-PEER PARADIGM

Figure 1.59 Centralized network



PEER-TO-PEER PARADIGM

Centralized Networks

- A **peer**, looking for a particular file, sends a **query** to a central server.
- The **server** searches its directory and **responds with the IP addresses of nodes** that have a copy of the file.
- The peer contacts one of the nodes and downloads the file.
- The directory is **constantly updated** as nodes join or leave the peer.
- Centralized networks make the maintenance of the directory simple but have several **drawbacks**.
- Accessing the directory can generate **huge traffic** and **slow down** the system.
- The central servers are **vulnerable to attack**, and if all of them fail, the whole system goes down.

PEER-TO-PEER PARADIGM

Decentralized Network

- A decentralized P2P network **does not depend on a centralized directory system.**
- In this model, peers arrange themselves into an **overlay network**, which is a **logical network** made on top of the physical network.
- Depending on **how the nodes in the overlay network are linked**, a decentralized P2P network is classified as either **unstructured or structured.**

PEER-TO-PEER PARADIGM

Unstructured Networks

- In an unstructured P2P network, the nodes are linked **randomly**.
- The **Gnutella** network is an example of a peer-to-peer network that is **decentralized but unstructured**.

Structured Networks

- A structured network uses **a predefined set of rules to link nodes** so that a query can be effectively and efficiently resolved.
- The most common technique used for this purpose is the **Distributed Hash Table (DHT)**.
- DHT is used in many applications including Distributed Data Structure (**DDS**), Content Distributed Systems (**CDS**), Domain Name System (**DNS**), and **P2P file sharing**.
- One popular P2P file sharing protocol that uses the DHT is **BitTorrent**.

PEER-TO-PEER PARADIGM

Distributed Hash Table (DHT)

- A Distributed Hash Table (DHT) distributes data (or references to data) among a set of nodes according to some predefined rules.
- Each peer in a DHT-based network becomes responsible for a range of data items.
- To avoid the flooding overhead, DHT-based networks allow each peer to have a partial knowledge about the whole network.
- This knowledge can be used to route the queries about the data items to the responsible nodes using effective and scalable procedures.
- There are several protocols that implement DHT systems: Chord, Pastry, and Kademlia.

PEER-TO-PEER PARADIGM

Distributed Hash Table (DHT)


Address Space

- In a DHT-based network, each data item and the peer is mapped to a point in a large address of size 2^m .
- The address space is designed using modular arithmetic, which means that the points in the address space are distributed evenly on a circle with 2^m points (0 to $2^m - 1$) using clockwise direction as shown in Figure 1.60.
- Most of the DHT implementations use $m = 160$.

Figure 1.60 Address space

Note:

1. Space range is 0 to $2^m - 1$.
2. Calculation is done modulo 2^m .



PEER-TO-PEER PARADIGM

Distributed Hash Table (DHT)

Hashing Peer Identifier

- The first step in creating the DHT system is to place all peers on the address space **ring**. This is normally done by using **a hash function** that hashes the peer identifier, normally its IP address, to an m-bit integer, called a node ID.
node ID = hash (Peer IP address)

Hashing Object Identifier

- The **name of the object** (for example, a file) to be shared is also hashed to an m-bit integer in the same address space. The result in DHT is called **a key**.
key = hash (Object name)

PEER-TO-PEER PARADIGM

Chord

- Chord was published by Stoica et al in 2001.

Identifier Space

- Data items and nodes in Chord are **m-bit identifiers** that create an identifier space of size 2^m points distributed in a circle in the clockwise direction.
- We refer to the **identifier of a data item** as k (for key) and the **identifier of a peer** as N (for node).

Finger Table

- A node in the Chord algorithm **should be able to resolve a query**: given a key, the node should be able to find the **node identifier responsible for that key or forward** the query to another node.
- Forwarding means that each node needs to have a **routing table**. Chord requires that each node knows about m successor nodes and one predecessor node. Each node creates a routing table, called a **finger table** by Chord.

PEER-TO-PEER PARADIGM

Chord

Table 1.9 *Finger table*

i	<i>Target Key</i>	<i>Successor of Target Key</i>	<i>Information about Successor</i>
1	$N + 1$	Successor of $N + 1$	IP address and port of successor
2	$N + 2$	Successor of $N + 2$	IP address and port of successor
\vdots	\vdots	\vdots	\vdots
m	$N + 2^{m-1}$	Successor of $N + 2^{m-1}$	IP address and port of successor

[Behrouz A Forouzan, Firouz Mosharraf, "Computer Networks: A top down Approach", McGraw Hill Education]

PEER-TO-PEER PARADIGM

Chord

Interface

- Chord needs a set of operations referred to as the **Chord interface**.

Lookup

- Probably the **mostly used operation** in Chord is the **lookup**.
- Chord is designed to let **peers share available services** between themselves.
- **To find the object to be shared**, a peer needs to know the node that is responsible for that object: **the peer that stores a reference to that object**.
- In Chord, **a peer that is the successor of a set of keys in the ring is the responsible peer for those keys**.
- **Finding the responsible node** is actually finding the successor of a key.

PEER-TO-PEER PARADIGM

A Popular P2P Network: BitTorrent

- BitTorrent is a P2P protocol, designed by Bram Cohen, for sharing a large file among a set of peers.
- However, the term sharing in this context is different from other filesharing protocols.
- Instead of one peer allowing another peer to download the whole file, a group of peers takes part in the process to give all peers in the group a copy of the file.
- File sharing is done in a collaborating process called a torrent.

PEER-TO-PEER PARADIGM

A Popular P2P Network: BitTorrent

- Each peer participating in a torrent downloads chunks of the large file from another peer that has it and uploads chunks of that file to other peers that do not have it, a kind of **tit-for-tat**, a trading game played by kids.
- The set of all peers that takes part in a torrent is referred to as a **swarm**.
- A peer in a swarm that has the complete content file is called a **seed**; a peer that has only part of the file and wants to download the rest is called a **leech**.
- A swarm is a **combination of seeds and leeches**.
- BitTorrent has gone through several versions and implementations, the original one which uses a central node called a **tracker**.