

Routing Algorithms

- Typically a host is attached directly to one router, the **default router** for the host (also called the **first-hop router** for the host).
- Whenever a host sends a packet, the packet is transferred to its **default router**.
- We refer to the default router of the source host as the **source router** and the default router of the destination host as the **destination router**.

Routing Algorithms

- The purpose of a routing algorithm is simple: given a set of routers, with links connecting the routers, a routing algorithm finds a “good” path from source router to destination router.
- Typically, a good path is one that has the least cost

Routing Algorithms

- A graph is used to formulate routing problems.
- The **graph** $G = (N, E)$ is a set N of nodes and a collection E of edges, where each edge is a pair of nodes from N .
- In the context of network-layer routing, the nodes in the graph represent **routers** - the points at which packet forwarding decisions are made - and the edges connecting these nodes represent **the physical links** between these routers.
- Such a graph is **an abstraction of a computer network**.

Routing Algorithms





Figure 4.27 ♦ Abstract graph model of a computer network

graph: $G = (N, E)$

N = set of routers = $\{ u, v, w, x, y, z \}$

E = set of links = $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

Routing Algorithms



cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

key question: what is the least-cost path between u and z?

routing algorithm: algorithm that finds that least cost path

Graph Abstraction


- An **edge also has a value** representing its cost.
- An edge's cost reflect the **physical length** of the corresponding link, the **link speed**, or the **monetary cost** associated with a link.
- For any edge (x,y) in E , we denote $c(x,y)$ as the cost of the edge between nodes x and y .
- If the pair (x,y) **does not belong to E** , we set $c(x,y) = \infty$.
- Consider only **undirected graphs** (i.e., graphs whose edges do not have a direction), so that edge (x,y) is the same as edge (y,x) and that **$c(x,y) = c(y,x)$** .
- Also, a node y is said to be a **neighbor** of node x if (x,y) belongs to E .

Graph Abstraction

- A natural **goal of a routing algorithm** is to identify the **least costly paths** between sources and destinations.
- To make this problem more precise, recall that a **path** in a graph $G = (N, E)$ is a sequence of nodes (x_1, x_2, \dots, x_p) such that each of the pairs $(x_1, x_2), (x_2, x_3), \dots, (x_{p-1}, x_p)$ are **edges** in E .
- The **cost** of a path (x_1, x_2, \dots, x_p) is simply the **sum of all the edge costs** along the path, that is, $c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$.

Graph Abstraction

- Given any two nodes x and y , there are typically many paths between the two nodes, with each path having a cost. One or more of these paths is a **least-cost path**.
- The **least-cost problem** is therefore clear: **Find a path between the source and destination that has least cost.**
- For example, the least-cost path between source node u and destination node w is (u, x, y, w) with **a path cost of 3.**
- Note that if all edges in the graph have the same cost, **the least-cost path is also the shortest path.**



Classification of Routing Algorithms

1. Global routing algorithm / Decentralized routing algorithm
2. Static routing algorithms / Dynamic routing algorithms
3. Load-sensitive algorithm / Load-insensitive algorithm

Global routing algorithm

- A **global routing algorithm** computes the least-cost path between a source and destination using complete, **global knowledge** about the network.
- That is, the algorithm takes the connectivity between all nodes and all link costs as inputs.
- This then requires that the algorithm somehow obtain this information before actually performing the calculation.
- The calculation itself can be run at one site.
- In practice, algorithms with global state information are often referred to as **link-state (LS) algorithms**, since the algorithm must be aware of the cost of each link in the network.

Decentralized routing algorithm

- In a **decentralized routing algorithm**, the calculation of the least-cost path is carried out in an **iterative, distributed manner**.
- **No node** has complete information about the costs of all network links.
- Instead, **each node begins with only the knowledge of the costs of its own directly attached links**.
- Then, through an **iterative process** of calculation and exchange of information with its neighboring nodes (that is, nodes that are at the other end of links to which it itself is attached), a node gradually calculates the least-cost path to a destination or set of destinations.
- The decentralized routing algorithm is called a **distance-vector (DV)** algorithm, because each node maintains a vector of estimates of the costs (distances) to all other nodes in the network.

Dynamic routing algorithms and Static routing algorithms

- In **Static routing algorithms**, routes change very slowly over time, often as a result of human intervention (for example, a human manually editing a router's forwarding table).
- **Dynamic routing algorithms** change the routing paths as the network traffic loads or topology change.
- A dynamic algorithm can be run **either periodically or in direct response to topology or link cost changes**.
- While dynamic algorithms are **more responsive** to network changes, they are also **more susceptible to problems** such as routing loops and oscillation in routes.

Load-sensitive algorithm / Load-insensitive algorithm

- In a **load-sensitive algorithm**, link costs vary dynamically to reflect the **current level of congestion** in the underlying link.
- If a high cost is associated with a link that is currently congested, a routing algorithm will tend to choose routes around such a congested link.
- Today's Internet routing algorithms (such as RIP, OSPF, and BGP) are **load-insensitive**, as a link's cost does not explicitly reflect its current (or recent past) level of congestion.

The Link-State (LS) Routing Algorithm

- The link-state routing algorithm is known as Dijkstra's algorithm, named after its inventor.
- A closely related algorithm is Prim's algorithm;
- Dijkstra's algorithm computes the least-cost path from one node (the source, which we will refer to as u) to all other nodes in the network.
- Dijkstra's algorithm is iterative and has the property that after the k^{th} iteration of the algorithm, the least-cost paths are known to k destination nodes, and among the least-cost paths to all destination nodes, these k paths will have the k smallest costs.

The Link-State (LS) Routing Algorithm

Let us define the following notation:

- $D(v)$: **cost of the least-cost path** from the source node to **destination v** as of this iteration of the algorithm.
- $p(v)$: **previous node** (neighbor of v) along the current least-cost path from the source to v .
- N' : **subset of nodes**; v is in N' if the least-cost path from the source to v is definitively know

The Link-State (LS) Routing Algorithm

Link-State (LS) Algorithm for Source Node u

[James F Kurose and Keith W Ross, “Computer Networking: A Top - Down Approach”, Pearson Education; 6 th Edition (2017)]


```
1  Initialization:
2       $N' = \{u\}$ 
3      for all nodes  $v$ 
4          if  $v$  is a neighbor of  $u$ 
5              then  $D(v) = c(u,v)$ 
6              else  $D(v) = \infty$ 
7
8  Loop
9      find  $w$  not in  $N'$  such that  $D(w)$  is a minimum
10     add  $w$  to  $N'$ 
11     update  $D(v)$  for each neighbor  $v$  of  $w$  and not in  $N'$ :
12          $D(v) = \min( D(v), D(w) + c(w,v) )$ 
13     /* new cost to  $v$  is either old cost to  $v$  or known
14        least path cost to  $w$  plus cost from  $w$  to  $v$  */
15 until  $N' = N$ 
```

- The **number of times** the loop is executed is equal to the **number of nodes** in the network.
- The algorithm will have calculated the **shortest paths** from the source node u to every other node in the network.

The Link-State (LS) Routing Algorithm

step	N'	$D(v), p(v)$	$D(w), p(w)$	$D(x), p(x)$	$D(y), p(y)$	$D(z), p(z)$
0	u	2,u	5,u	1,u	∞	∞
1	ux	2,u	4,x		2,x	∞
2	uxy	2,u	3,y			4,y
3	uxyv		3,y			4,y
4	uxyvw					4,y
5	uxyvwz					

Table 4.3 ♦ Running the link-state algorithm on the network



The Link-State (LS) Routing Algorithm

- In the **initialization step**, the currently known least-cost paths from u to its directly attached neighbors, v , x , and w , are initialized to 2, 1, and 5, respectively.
- Note in particular that the cost to **w is set to 5** (even though we will soon see that a lesser-cost path does indeed exist) since this is the **cost of the direct (one hop) link from u to w** .
- The costs to y and z are set to **infinity** because they are not directly connected to u .

The Link-State (LS) Routing Algorithm



- In the first iteration, we look among those nodes **not yet added to the set N'** and find that node **with the least cost** as of the end of the previous iteration.
- That node is x , with a cost of 1, and thus x is added to the set N' .
- Line 12 of the LS algorithm is then performed to update $D(v)$ for all nodes v , yielding the results shown in the second line (Step 1) in Table 4.3.
- The cost of the path to v **is unchanged**.
- The cost of the path to w (which was 5 at the end of the initialization) through **node x is found to have a cost of 4**. Hence this **lower-cost path is selected** and w 's predecessor along the shortest path from u is set to x .
- Similarly, the cost to y (through x) is computed to be 2, and the table is **updated accordingly**.

The Link-State (LS) Routing Algorithm

- In the second iteration, nodes v and y are found to have the **least-cost paths (2)**, and we break the tie arbitrarily and add y to the set N so that N now contains u , x , and y .
- The cost to the remaining nodes not yet in N' , that is, nodes v , w , and z , are updated via line 12 of the LS algorithm, yielding the results.

The algorithm calculated the **shortest paths from the source node u to every other node in the network.**

The Link-State (LS) Routing Algorithm



Destination	Link
v	(u, v)
w	(u, x)
x	(u, x)
y	(u, x)
z	(u, x)

Figure 4.28 ♦ Least cost path and forwarding table for node u

[James F Kurose and Keith W Ross, "Computer Networking: A Top - Down Approach", Pearson Education; 6 th Edition (2017)]

The algorithm calculated the **shortest paths from the source node u to every other node in the network.**

The Link-State (LS) Routing Algorithm

- What is the **computational complexity** of this algorithm? That is, given n nodes (not counting the source), how much computation must be done in the worst case to find the least-cost paths from the source to all destinations?

The preceding implementation of the LS algorithm has worst-case complexity of **order n squared: $O(n^2)$** .

The Link-State (LS) Routing Algorithm

- In the first iteration, we need to **search through all n nodes** to determine the node, w , not in N' that has the minimum cost.
- In the second iteration, we need **to check $n - 1$ nodes** to determine the minimum cost; in the third iteration $n - 2$ nodes, and so on.
- Overall, the total number of nodes we need to search through over all the iterations is **$n(n + 1)/2$** .

The preceding implementation of the LS algorithm has worst-case complexity of order n squared: **$O(n^2)$** .

The Distance-Vector (DV) Routing Algorithm

In distance vector routing, each node shares its **routing table** with its immediate neighbors **periodically** and when there is a change.

The Distance-Vector (DV) Routing Algorithm

- LS algorithm is an algorithm using global information, the Distance Vector (DV) algorithm is iterative, asynchronous, and distributed.
- It is distributed in that each node receives some information from one or more of its directly attached neighbors, performs a calculation, and then distributes the results of its calculation back to its neighbors.
- It is iterative in that this process continues on until no more information is exchanged between neighbors.
- The algorithm is also self-terminating—there is no signal that the computation should stop; it just stops.

The Distance-Vector (DV) Routing Algorithm

- The algorithm is **asynchronous** in that it **does not require all of the nodes to operate** in lockstep with each other.
- An asynchronous, iterative, self-terminating, distributed algorithm is much more interesting and fun **than a centralized algorithm!**

Distance Vector Algorithm

Bellman-Ford Equation (dynamic programming)

Define

$d_x(y) :=$ cost of least-cost path from x to y

Then

$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$

where \min_v is taken **over all neighbors of x** .


- **Intuition:** After traveling from x to v , if we then take the least-cost path from v to y , the path cost will be $c(x,v) + d_v(y)$.
- Since **we must begin by traveling to some neighbor v** , the **least cost** from x to y is the minimum of $c(x,v) + d_v(y)$ taken over all neighbors v

Bellman-Ford example

Clearly, $d_A(Y) = 3$, $d_B(Y) = 7$

B-F equation says:

$$\begin{aligned} d_X(Y) &= \min \{ c(X,A) + d_A(Y), \\ &\quad c(X,B) + d_B(Y) \} \\ &= \min \{ 5 + 3, 6 + 7 \} = 8 \end{aligned}$$



Important practical contribution of the **Bellman-Ford equation** is that it suggests the form of the **neighbor-to-neighbor communication** that will take place in the DV algorithm

Distance Vector Algorithm

With the DV algorithm, each node x maintains the following routing data.

1. For each neighbor v , the cost $c(x,v)$ from x to directly attached neighbor v .

2. Node x 's distance vector. Distance vector: $\mathbf{D}_x = [D_x(y): y \in N]$ containing x 's estimate of its cost to all destinations y in N .

3. The distance vectors of each of its neighbors .

$$\mathbf{D}_v = [D_v(y): y \in N]$$

for each neighbor v of x .

Distance vector algorithm

Basic idea:

- Each node **periodically sends its own distance vector** estimate to neighbors
- When a node **x receives new DV estimate** from neighbor, it **updates its own DV** using **B-F equation**:

$$D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\} \quad \text{for each node } y \in N$$

Distance vector algorithm

Iterative each local iteration caused by:

- local link cost change
- DV update message from neighbor.
- Continue on until no more info is exchanged between neighbors.

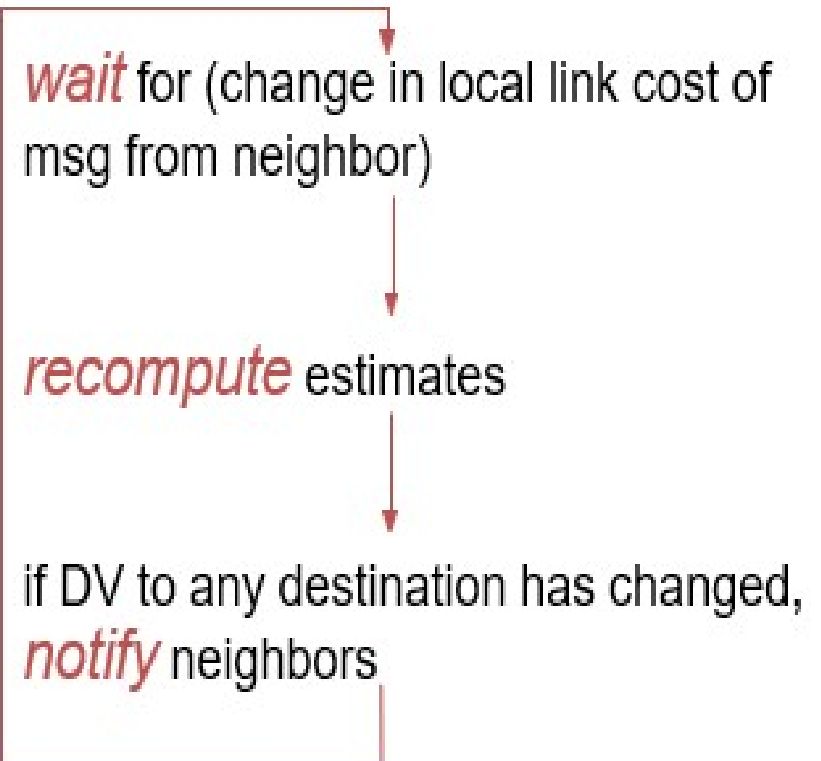
Distributed:

- Each node receives some info from one or more of its directly attached neighbors.
- Performs a calculation and then distributes the results of its calculation back to its neighbors.
- each node notifies neighbors *only* when its DV changes
 - neighbors then notify their neighbors if necessary

Self Terminating:

Asynchronous: it doesn't require all of the nodes to operate in lockstep with each other

Each node:



Distance vector algorithm

Working:

- Each node x begins with $D_x(y)$, an estimate of the cost of the least-cost path from itself to node y , for all nodes in N .
- $D_x = [D_x(y): y \text{ in } N]$ be node x 's **distance vector**, which is the vector of cost estimates from x to all other nodes, y in N .
- Each node x maintains the following routing information:
 - For each neighbor v , the cost $c(x,v)$ from x to directly attached neighbor, v
 - Node x 's distance vector, that is, $D_x = [D_x(y): y \text{ in } N]$, containing x 's estimate of its cost to all destinations, y in N
 - The distance vectors of each of its neighbors, that is, $D_v = [D_v(y): y \text{ in } N]$ for each neighbor v of x .

Distance vector algorithm

- From time to time, each node sends a copy of its distance vector to each of its neighbors.
- When a node x receives a new distance vector from any of its neighbors v , it saves v 's distance vector, and then uses the Bellman-Ford equation to update its own distance vector as follows:

$$D_x(y) = \min_v \{c(x,v) + D_v(y)\} \text{ for each node } y \text{ in } N$$

- If node x 's distance vector has changed as a result of this update step, node x will then send its updated distance vector to each of its neighbors, which can in turn update their own distance vectors.

Distance vector algorithm

$$D_x(x)=0$$

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

node x table

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞

node y table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

node z table

	cost to	x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0


	cost to	x	y	z
from	x	0	2	7
	y	2	0	1
	z	7	1	0

	cost to	x	y	z
x		0	2	7
y		2	0	1
z		3	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

	cost to	x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0



Distance vector algorithm

- At initialization node x has not received anything from node y or z, the entries in the second and third rows are **initialized to infinity**

$$D_z(x) = \min(C(z, x) + D_x(x), C(z, y) + D_y(x)) \\ = \min(7 + 0, 1 + 2) = 3$$

$$D_y(x) = \min(C(y, x) + D_x(x), C(y, z) + D_z(x)) \\ = \min(2 + 0, 1 + 3) = 2$$

Distance vector algorithm

Distance-Vector (DV) Algorithm

At each node, x :

```
1  Initialization:
2    for all destinations  $y$  in  $N$ :
3       $D_x(y) = c(x,y)$  /* if  $y$  is not a neighbor then  $c(x,y) = \infty$  */
4    for each neighbor  $w$ 
5       $D_w(y) = ?$  for all destinations  $y$  in  $N$ 
6    for each neighbor  $w$ 
7      send distance vector  $D_x = [D_x(y): y \text{ in } N]$  to  $w$ 
8
9  loop
10   wait (until I see a link cost change to some neighbor  $w$  or
11         until I receive a distance vector from some neighbor  $w$ )
12
13   for each  $y$  in  $N$ :
14      $D_x(y) = \min_v \{c(x,v) + D_v(y)\}$ 
15
16   if  $D_x(y)$  changed for any destination  $y$ 
17     send distance vector  $D_x = [D_x(y): y \text{ in } N]$  to all neighbors
18
19 forever
```